


# RNN 실습 - 텐서플로우에서 LSTM 및 GRU 사용

1. RNN 모형 및 구성방법
2. 텐서플로우에서 지원하는 RNN 'Cell' 유형
3. 유형별 특성 테스트 위한 코드 구성
4. Test #1 - Vanila RNN
5. Test #2 - Basic LSTM
6. Test #3 - GRU
7. Test #4 - LSTMCell + forget\_bias
8. Test #5 - LayerNormBasicLSTMCell
9. Test #6 - LayerNormBasicLSTMCell - what's wrong?
10. 정리

```
In [1]: !rm -fr logdir
        !mkdir -p logdir
```

```
In [2]: %load_ext do_not_print_href
        %matplotlib inline
        from __future__ import print_function, division
        import sys
        import time
        import numpy as np
        import tensorflow as tf
        import matplotlib.pyplot as plt
```

## RNN 모형 (복습)



이미지 출처: <http://www.wildml.com/2015/09/recurrent-neural-networks-tutorial-part-1-introduction-to-rnns/>

$$s_t = \tanh(Ux_t + Ws_{t-1})$$

$$o_t = \text{softmax}(Vs_t)$$

하지만 텐서플로우 RNN 은...



- $V$ 에 해당하는 구조가 없음
- **softmax()** 도 없음
- $V$ 와 **softmax()** 는 필요할 때만 만들어서 붙이면 됨 (`tf.layers.dense`, `tf.nn.softmax`)

## RNN 모델 구성 (복습)

- [tf.contrib.rnn.BasicRNNCell](#)

```
`__init__`(  
    num_units,  
    activation=None,  
    reuse=None  
)
```

- [tf.nn.dynamic\\_rnn](#)

```
dynamic_rnn(  
    cell,  
    inputs,  
    sequence_length=None,  
    initial_state=None,  
    dtype=None,  
    parallel_iterations=None,  
    swap_memory=False,  
    time_major=False,  
    scope=None  
)
```

- RNN 구성

```
cell                = tf.contrib.rnn.BasicRNNCell(  
                        num_hidden_units)  
last, states        = tf.nn.dynamic_rnn(  
                        cell,  
                        inputs,  
                        sequence_length=sequence_length,  
                        dtype=tf.float32)
```

## RNN Cell 종류

- [tf.contrib.rnn.BasicRNNCell](#)

...

- [tf.contrib.rnn.BasicLSTMCell](#)

```

`__init__`(
    num_units,
    forget_bias=1.0,
    state_is_tuple=True,
    activation=None,
    reuse=None
)

```

- [tf.contrib.rnn.LSTMCell](#)

```

`__init__`(
    num_units,
    use_peepholes=False,
    cell_clip=None,
    initializer=None,
    num_proj=None,
    proj_clip=None,
    num_unit_shards=None,
    num_proj_shards=None,
    forget_bias=1.0,
    state_is_tuple=True,
    activation=None,
    reuse=None
)

```

- [tf.contrib.rnn.GRUCell](#)

```

`__init__`(
    num_units,
    activation=None,
    reuse=None,
    kernel_initializer=None,
    bias_initializer=None
)

```

- [tf.contrib.rnn.LayerNormBasicLSTMCell](#)

LSTM unit with **layer normalization** and **recurrent dropout**.

- "Layer Normalization" Jimmy Lei Ba, Jamie Ryan Kiros, Geoffrey E. Hinton
  - <https://arxiv.org/abs/1607.06450>
- "Recurrent Dropout without Memory Loss" Stanislau Semeniuta, Aliaksei Severyn, Erhardt Barth.
  - <https://arxiv.org/abs/1603.05118>
- "Layer Normalization, Nishida Geio, 번역 김홍배 (slideshare.net)"
  - <https://www.slideshare.net/ssuser06e0c5/normalization-72539464>

```

`__init__`(
    num_units,
    forget_bias=1.0,
    input_size=None,
    activation=tf.tanh,
    layer_norm=True,
    norm_gain=1.0,
    norm_shift=0.0,
    dropout_keep_prob=1.0,
    dropout_prob_seed=None,
    reuse=None
)

```

- 기타 여러가지 Wrapper 지원 - [https://www.tensorflow.org/api\\_guides/python/contrib.rnn](https://www.tensorflow.org/api_guides/python/contrib.rnn)

```

In [3]: from tensorflow.examples.tutorials.mnist.input_data \
        import read_data_sets

```

```

mnist = read_data_sets('./mnist', one_hot=False)

```

```

Extracting ./mnist/train-images-idx3-ubyte.gz
Extracting ./mnist/train-labels-idx1-ubyte.gz
Extracting ./mnist/t10k-images-idx3-ubyte.gz
Extracting ./mnist/t10k-labels-idx1-ubyte.gz

```

```

In [4]: INPUT_UNITS = 28
        NUM_HIDDEN_UNITS = 31

        BATCH_SIZE = 128
        MAX_SEQ_LEN = 28

```

```

In [5]: class MnistRnn:
        def __init__(self,
                      inputs,
                      labels,
                      input_units,
                      num_hidden_units,
                      batch_size,
                      max_seq_len,
                      rnn_cell_class = tf.contrib.rnn.BasicRNNCell,
                      # 여기 이후의 인수는 잠깐 무시 해 주세요
                      add_check = False,
                      lr = 0.001,
                      use_grad_clip = False):
            '''
            inputs: in shape [batch_size, max_seq_len, input_size]
            labels: in shape [batch_size]
            '''

            cell = rnn_cell_class(num_hidden_units)
            sequence_length = [max_seq_len] * batch_size
            last, states = tf.nn.dynamic_rnn(
                cell,

```

```

        inputs,
        sequence_length=sequence_length,
        dtype=tf.float32)

# max_seq_len 축으로 0~27 까지 값 중에
# 0~26 때의 출력 값은 사용하지 않음
rnn_output = last[:,max_seq_len-1,:]
# outputs shape will be: [batch_size, 10]
outputs     = tf.layers.dense(rnn_output, 10)
loss        = tf.losses.sparse_softmax_cross_entropy(
                labels, outputs)

if use_grad_clip:
    tvars_    = tf.trainable_variables()
    grads_, _ = tf.clip_by_global_norm(
                    tf.gradients(
                        loss,
                        tvars_),
                    5.0)
    optimize  = \
        tf.train.AdamOptimizer(learning_rate=lr). \
            apply_gradients(zip(grads_, tvars_))
else:
    optimize  = \
        tf.train.AdamOptimizer(learning_rate=lr). \
            minimize(loss)

# accuracy
preds        = tf.argmax(outputs, axis=1)
errors       = tf.count_nonzero(labels - preds)
accuracy     = 1.0 - tf.cast(errors,tf.float32) / \
                tf.cast(tf.size(preds),tf.float32)

# 클래스 객체 외부에서 참고할 수 있도록 속성으로 저장
self.outputs    = outputs
self.loss       = loss
self.optimize   = optimize
self.accuracy   = accuracy

# check_numerics
self.check = [tf.check_numerics(
                    t,
                    'check_numerics: {}'.format(t.name)) \
                for t in tf.gradients(
                    loss,
                    tf.trainable_variables()) \
                if t is not None] \
    if add_check \
    else tf.constant(1.0)

```

```

In [6]: train_loop_count = mnist.train.num_examples // BATCH_SIZE
        test_loop_count  = mnist.test.num_examples // BATCH_SIZE

        train_loop_count, test_loop_count

```

```

Out[6]: (429, 78)

```

```

In [7]: def train(inputs, labels, max_epochs, train_writer, test_writer):
        step = 0
        for ep in range(max_epochs):

            train_elapsed = []
            train_losses = []
            train_accuracy = []
            for i in range(train_loop_count):
                t_start = time.time()
                offs = i * BATCH_SIZE
                batch_input = \
                    mnist.train.images[offs:offs+BATCH_SIZE,:]
                batch_input = \
                    batch_input.reshape(
                        [BATCH_SIZE,
                         MAX_SEQ_LEN,
                         INPUT_UNITS])
                batch_label = \
                    mnist.train.labels[offs:offs+BATCH_SIZE]
                optimize, loss, accuracy, _ = \
                    sess.run([model.optimize,
                              model.loss,
                              model.accuracy,
                              model.check],
                              feed_dict = {
                                  inputs: batch_input,
                                  labels: batch_label })
                train_losses.append(loss)
                train_accuracy.append(accuracy)
                t_elapsed = time.time() - t_start
                train_elapsed.append(t_elapsed)

            step += 1
            summary = tf.Summary(
                value=[
                    tf.Summary.Value(
                        tag='train_accuracy',
                        simple_value=accuracy
                    ),
                    tf.Summary.Value(
                        tag='loss',
                        simple_value=loss
                    ),
                ]
            )
            train_writer.add_summary(summary, global_step=step)

            if step % 250 == 0:
                print(('[trn] ep {:d}, step {:d}, ' +
                      'loss {:f}, accu {:f}, ' +
                      'sec/iter {:f}').format(
                        ep + 1,
                        step,
                        np.mean(train_losses),
                        np.amin(train_accuracy),
                        np.mean(train_elapsed)))
                train_losses = []

```

```

        train_accuracy = []
        train_elapsed = []

    train_writer.flush()

    test_elapsed = []
    test_accuracy = []

    for i in range(test_loop_count):
        t_start = time.time()
        offs = i * BATCH_SIZE
        batch_input = mnist.test.images[offs:offs+BATCH_SIZE,:]
        batch_input = batch_input.reshape(
            [BATCH_SIZE,
             MAX_SEQ_LEN,
             INPUT_UNITS])
        batch_label = mnist.test.labels[offs:offs+BATCH_SIZE]
        accuracy, = \
            sess.run([model.accuracy],
                     feed_dict = {
                         inputs: batch_input,
                         labels: batch_label })
        test_accuracy.append(accuracy)
        t_elapsed = time.time() - t_start
        test_elapsed.append(t_elapsed)

        step += 1

    if len(test_accuracy) > 0:
        print(('[tst] ep {:d}, step {:d}, ' +
              'accu {:f}, sec/iter {:f}').format(
            ep + 1,
            step,
            np.amin(test_accuracy),
            np.mean(test_elapsed)))

        summary = tf.Summary(
            value=[
                tf.Summary.Value(
                    tag='test_accuracy',
                    simple_value=np.amin(test_accuracy)
                ),
            ]
        )
        test_writer.add_summary(summary, global_step=step)
        test_writer.flush()

```

## Test #1 - Vanila RNN

```
cell = tf.contrib.rnn.BasicRNNCell(num_hidden_units)
```



```
In [8]: tf.reset_default_graph()

inputs_ = tf.placeholder(tf.float32,
                          [BATCH_SIZE, MAX_SEQ_LEN, INPUT_UNITS],
                          name='inputs')
labels_ = tf.placeholder(tf.int64,
                          [BATCH_SIZE],
                          name='labels')

model = MnistRnn(inputs_,
                  labels_,
                  INPUT_UNITS,
                  NUM_HIDDEN_UNITS,
                  BATCH_SIZE,
                  MAX_SEQ_LEN,
                  tf.contrib.rnn.BasicRNNCell)
```

```
In [9]: config = tf.ConfigProto(gpu_options={'allow_growth':True})
sess = tf.InteractiveSession(config=config)

tf.global_variables_initializer().run()

train_writer = tf.summary.FileWriter('logdir/train_basic_rnn',
                                     graph=tf.get_default_graph())
test_writer  = tf.summary.FileWriter('logdir/test_basic_rnn',
                                     graph=tf.get_default_graph())

train(inputs_, labels_, 10, train_writer, test_writer)
```

```
[trn] ep 1, step 250, loss 1.550250, accu 0.085938, sec/iter 0.007951
[tst] ep 1, step 507, accu 0.468750, sec/iter 0.002797
[trn] ep 2, step 750, loss 0.987132, accu 0.562500, sec/iter 0.007306
[tst] ep 2, step 1014, accu 0.539062, sec/iter 0.002285
[trn] ep 3, step 1250, loss 0.765099, accu 0.625000, sec/iter 0.007807
[tst] ep 3, step 1521, accu 0.648438, sec/iter 0.004670
[trn] ep 4, step 1750, loss 0.632159, accu 0.648438, sec/iter 0.007768
[tst] ep 4, step 2028, accu 0.656250, sec/iter 0.003359
[trn] ep 5, step 2250, loss 0.562881, accu 0.664062, sec/iter 0.009445
[tst] ep 5, step 2535, accu 0.671875, sec/iter 0.005474
[trn] ep 6, step 2750, loss 0.504307, accu 0.695312, sec/iter 0.008914
[tst] ep 6, step 3042, accu 0.679688, sec/iter 0.003352
[trn] ep 7, step 3250, loss 0.462129, accu 0.718750, sec/iter 0.009098
[tst] ep 7, step 3549, accu 0.664062, sec/iter 0.005065
[trn] ep 8, step 3750, loss 0.429448, accu 0.757812, sec/iter 0.009314
[tst] ep 8, step 4056, accu 0.710938, sec/iter 0.003061
[trn] ep 9, step 4250, loss 0.401605, accu 0.750000, sec/iter 0.008644
[tst] ep 9, step 4563, accu 0.710938, sec/iter 0.004161
[trn] ep 10, step 4750, loss 0.384211, accu 0.781250, sec/iter 0.008817
[tst] ep 10, step 5070, accu 0.726562, sec/iter 0.003440
```

```
In [10]: # !tensorboard --ip 0.0.0.0 --logdir logdir
```

## Test #2 - Basic LSTM

```
cell = tf.contrib.rnn.BasicLSTMCell(num_hidden_units)
```

```
In [11]: tf.reset_default_graph()

inputs_ = tf.placeholder(tf.float32,
                          [BATCH_SIZE, MAX_SEQ_LEN, INPUT_UNITS],
                          name='inputs')
labels_ = tf.placeholder(tf.int64,
                          [BATCH_SIZE],
                          name='labels')

model = MnistRnn(inputs_,
                  labels_,
                  INPUT_UNITS,
                  NUM_HIDDEN_UNITS,
                  BATCH_SIZE,
                  MAX_SEQ_LEN,
                  tf.contrib.rnn.BasicLSTMCell)
```

```
In [12]: config = tf.ConfigProto(gpu_options={'allow_growth':True})
sess = tf.InteractiveSession(config=config)

tf.global_variables_initializer().run()

train_writer = tf.summary.FileWriter('logdir/train_basic_lstm',
                                     graph=tf.get_default_graph())
test_writer  = tf.summary.FileWriter('logdir/test_basic_lstm',
                                     graph=tf.get_default_graph())

train(inputs_, labels_, 10, train_writer, test_writer)
```

```
[trn] ep 1, step 250, loss 1.499541, accu 0.039062, sec/iter 0.013030
[tst] ep 1, step 507, accu 0.695312, sec/iter 0.005082
[trn] ep 2, step 750, loss 0.469044, accu 0.671875, sec/iter 0.013055
[tst] ep 2, step 1014, accu 0.796875, sec/iter 0.004687
[trn] ep 3, step 1250, loss 0.304381, accu 0.773438, sec/iter 0.013056
[tst] ep 3, step 1521, accu 0.851562, sec/iter 0.006747
[trn] ep 4, step 1750, loss 0.228537, accu 0.835938, sec/iter 0.013932
[tst] ep 4, step 2028, accu 0.867188, sec/iter 0.005018
[trn] ep 5, step 2250, loss 0.196097, accu 0.843750, sec/iter 0.012640
[tst] ep 5, step 2535, accu 0.867188, sec/iter 0.004851
[trn] ep 6, step 2750, loss 0.166952, accu 0.882812, sec/iter 0.012917
[tst] ep 6, step 3042, accu 0.875000, sec/iter 0.004824
[trn] ep 7, step 3250, loss 0.147182, accu 0.898438, sec/iter 0.014815
[tst] ep 7, step 3549, accu 0.867188, sec/iter 0.005175
[trn] ep 8, step 3750, loss 0.133158, accu 0.906250, sec/iter 0.013107
[tst] ep 8, step 4056, accu 0.867188, sec/iter 0.005694
[trn] ep 9, step 4250, loss 0.119845, accu 0.914062, sec/iter 0.013864
[tst] ep 9, step 4563, accu 0.906250, sec/iter 0.005666
[trn] ep 10, step 4750, loss 0.107291, accu 0.914062, sec/iter 0.015189
[tst] ep 10, step 5070, accu 0.906250, sec/iter 0.004921
```

```
In [13]: # !tensorboard --ip 0.0.0.0 --logdir logdir
```

## Test #3 - GRU

```
cell = tf.contrib.rnn.GRUCell(num_hidden_units)
```

```
In [14]: tf.reset_default_graph()

inputs_ = tf.placeholder(tf.float32,
                          [BATCH_SIZE, MAX_SEQ_LEN, INPUT_UNITS],
                          name='inputs')
labels_ = tf.placeholder(tf.int64,
                          [BATCH_SIZE],
                          name='labels')

model = MnistRnn(inputs_,
                  labels_,
                  INPUT_UNITS,
                  NUM_HIDDEN_UNITS,
                  BATCH_SIZE,
                  MAX_SEQ_LEN,
                  tf.contrib.rnn.GRUCell)
```

```
In [15]: config = tf.ConfigProto(gpu_options={'allow_growth':True})
sess = tf.InteractiveSession(config=config)

tf.global_variables_initializer().run()

train_writer = tf.summary.FileWriter(
    'logdir/train_gru',
    graph=tf.get_default_graph())
test_writer  = tf.summary.FileWriter(
    'logdir/test_gru',
    graph=tf.get_default_graph())

train(inputs_, labels_, 10, train_writer, test_writer)
```

```
[trn] ep 1, step 250, loss 1.550183, accu 0.117188, sec/iter 0.014488
[tst] ep 1, step 507, accu 0.679688, sec/iter 0.005059
[trn] ep 2, step 750, loss 0.464381, accu 0.703125, sec/iter 0.014379
[tst] ep 2, step 1014, accu 0.820312, sec/iter 0.004900
[trn] ep 3, step 1250, loss 0.284551, accu 0.820312, sec/iter 0.014348
[tst] ep 3, step 1521, accu 0.859375, sec/iter 0.004766
[trn] ep 4, step 1750, loss 0.216922, accu 0.828125, sec/iter 0.015113
[tst] ep 4, step 2028, accu 0.882812, sec/iter 0.004711
[trn] ep 5, step 2250, loss 0.180984, accu 0.851562, sec/iter 0.014021
[tst] ep 5, step 2535, accu 0.890625, sec/iter 0.005294
[trn] ep 6, step 2750, loss 0.158896, accu 0.835938, sec/iter 0.015809
[tst] ep 6, step 3042, accu 0.882812, sec/iter 0.004615
[trn] ep 7, step 3250, loss 0.140687, accu 0.875000, sec/iter 0.015920
[tst] ep 7, step 3549, accu 0.890625, sec/iter 0.007297
[trn] ep 8, step 3750, loss 0.126121, accu 0.898438, sec/iter 0.018148
[tst] ep 8, step 4056, accu 0.906250, sec/iter 0.005170
[trn] ep 9, step 4250, loss 0.115712, accu 0.914062, sec/iter 0.013839
[tst] ep 9, step 4563, accu 0.906250, sec/iter 0.004460
[trn] ep 10, step 4750, loss 0.106273, accu 0.914062, sec/iter 0.013993
[tst] ep 10, step 5070, accu 0.906250, sec/iter 0.004573
```

```
In [16]: # !tensorboard --ip 0.0.0.0 --logdir logdir
```

## Test #4 - LSTMCell + forget\_bias

```
cell = tf.contrib.rnn.BasicLSTMCell(
    num_hidden_units,
    forget_bias=lstm_forget_bias)
```

```
In [17]: tf.reset_default_graph()

inputs_ = tf.placeholder(tf.float32,
                          [BATCH_SIZE, MAX_SEQ_LEN, INPUT_UNITS],
                          name='inputs')
labels_ = tf.placeholder(tf.int64,
                          [BATCH_SIZE],
                          name='labels')

lstm_with_forget_bias = lambda num_hidden_units: \
    tf.contrib.rnn.BasicLSTMCell(
        num_hidden_units,
        forget_bias=5.0)

model = MnistRnn(inputs_,
                  labels_,
                  INPUT_UNITS,
                  NUM_HIDDEN_UNITS,
                  BATCH_SIZE,
                  MAX_SEQ_LEN,
                  lstm_with_forget_bias)
```

[illegible]

```
In [19]: train(inputs_, labels_, 10, train_writer, test_writer)

[trn] ep 1, step 250, loss 1.470503, accu 0.054688, sec/iter 0.013168
[tst] ep 1, step 507, accu 0.617188, sec/iter 0.005363
[trn] ep 2, step 750, loss 0.543548, accu 0.664062, sec/iter 0.013188
[tst] ep 2, step 1014, accu 0.703125, sec/iter 0.007593
[trn] ep 3, step 1250, loss 0.369156, accu 0.757812, sec/iter 0.015381
[tst] ep 3, step 1521, accu 0.796875, sec/iter 0.005513
[trn] ep 4, step 1750, loss 0.286988, accu 0.828125, sec/iter 0.013617
[tst] ep 4, step 2028, accu 0.820312, sec/iter 0.005127
[trn] ep 5, step 2250, loss 0.240008, accu 0.851562, sec/iter 0.015156
[tst] ep 5, step 2535, accu 0.835938, sec/iter 0.005836
[trn] ep 6, step 2750, loss 0.208887, accu 0.867188, sec/iter 0.013542
[tst] ep 6, step 3042, accu 0.859375, sec/iter 0.005498
[trn] ep 7, step 3250, loss 0.183400, accu 0.882812, sec/iter 0.015910
[tst] ep 7, step 3549, accu 0.851562, sec/iter 0.007850
[trn] ep 8, step 3750, loss 0.164586, accu 0.882812, sec/iter 0.014953
[tst] ep 8, step 4056, accu 0.867188, sec/iter 0.006749
[trn] ep 9, step 4250, loss 0.149524, accu 0.890625, sec/iter 0.015110
[tst] ep 9, step 4563, accu 0.882812, sec/iter 0.006931
[trn] ep 10, step 4750, loss 0.136078, accu 0.898438, sec/iter 0.015293
[tst] ep 10, step 5070, accu 0.898438, sec/iter 0.008359
```

```
In [20]: # !tensorboard --ip 0.0.0.0 --logdir logdir
```

## Test #5 - LayerNormBasicLSTMCell

```
cell = tf.contrib.rnn.LayerNormBasicLSTMCell(
    num_hidden_units)
```



```
In [21]: tf.reset_default_graph()

inputs_ = tf.placeholder(tf.float32,
                          [BATCH_SIZE, MAX_SEQ_LEN, INPUT_UNITS],
                          name='inputs')
labels_ = tf.placeholder(tf.int64,
                          [BATCH_SIZE],
                          name='labels')

model = MnistRnn(inputs_,
                  labels_,
                  INPUT_UNITS,
                  NUM_HIDDEN_UNITS,
                  BATCH_SIZE,
                  MAX_SEQ_LEN,
                  tf.contrib.rnn.LayerNormBasicLSTMCell)
```

```
In [22]: config = tf.ConfigProto(gpu_options={'allow_growth':True})
sess = tf.InteractiveSession(config=config)

tf.global_variables_initializer().run()

train_writer = tf.summary.FileWriter('logdir/train_ln_basic_lstm',
                                     graph=tf.get_default_graph())
test_writer  = tf.summary.FileWriter('logdir/test_ln_basic_lstm',
                                     graph=tf.get_default_graph())

train(inputs_, labels_, 1, train_writer, test_writer)

[trn] ep 1, step 250, loss nan, accu 0.039062, sec/iter 0.063241
[tst] ep 1, step 507, accu 0.054688, sec/iter 0.011702
```

```
In [23]: # !tensorboard --ip 0.0.0.0 --logdir logdir
```

## Test #6 - LayerNormBasicLSTMCell - what's wrong?

- `tf.check_numerics()`

When run, reports an `InvalidArgument` error if tensor has any values that are not a number (NaN) or infinity (Inf). Otherwise, passes tensor as-is.

```
check_numerics(
    tensor,
    message,
    name=None
)
```

- `tf.gradients()`

Constructs symbolic partial derivatives of sum of `ys` w.r.t. `x` in `xs`

```
gradients(
    ys,
    xs,
    grad_ys=None,
    name='gradients',
    colocate_gradients_with_ops=False,
    gate_gradients=False,
    aggregation_method=None
)
```

- 사용 예:

```
# check_numerics
self.check = [tf.check_numerics(t,
                                'check_numerics: {}'.format(t.name)) \
               for t in tf.gradients(
                   loss,
                   tf.trainable_variables()) \
               if t is not None]

...
summary, optimize, loss, accuracy, _ = \
    sess.run([model.train_summary,
              model.optimize,
              model.loss,
              model.accuracy,
              model.check],
             ...)
```

```
In [24]: tf.reset_default_graph()

inputs_ = tf.placeholder(tf.float32,
                          [BATCH_SIZE, MAX_SEQ_LEN, INPUT_UNITS],
                          name='inputs')
labels_ = tf.placeholder(tf.int64,
                          [BATCH_SIZE],
                          name='labels')

model = MnistRnn(inputs_,
                  labels_,
                  INPUT_UNITS,
                  NUM_HIDDEN_UNITS,
                  BATCH_SIZE,
                  MAX_SEQ_LEN,
                  tf.contrib.rnn.LayerNormBasicLSTMCell,
                  add_check = True)

# 여기서 NaN 문제가 생기는 걸 확인했으면, 다음 방법들을 시도 해 보세요
# 1. learning_rate 를 줄여
#     e.g.: lr = 0.0001
# 2. 적당한 값으로 gradient clipping
#     e.g.: use_grad_clip = True
# 3. 사용한 컴퍼넌트에 smoothing 할 수 있는 파라미터가 있는지 확인하고 적용
#     rnn_cell_class = lambda num_hidden_units: \
#         tf.contrib.rnn.LayerNormBasicLSTMCell(
#             num_hidden_units,
#             norm_gain=0.85,
#             norm_shift=0.15)
```

```
In [ ]: config = tf.ConfigProto(gpu_options={'allow_growth':True})
sess = tf.InteractiveSession(config=config)

tf.global_variables_initializer().run()

train_writer = tf.summary.FileWriter(
    'logdir/train_ln_basic_lstm_2',
    graph=tf.get_default_graph())
test_writer = tf.summary.FileWriter(
    'logdir/test_ln_basic_lstm_2',
    graph=tf.get_default_graph())

train(inputs_, labels_, 10, train_writer, test_writer)
```

```
"" InvalidArgumentError: check_numerics:
gradients_1/rnn/while/rnn/layer_norm_basic_lstm_cell/state_1/batchnorm/sub/Enter_grad/b_acc_3:0 :
Tensor had NaN values [[Node: CheckNumerics_10 = CheckNumerics[T=DT_FLOAT,
message="check_numerics:
gradients_1/rnn/while/rnn/layer_norm_basic_lstm_cell/state_1/batchnorm/sub/Enter_grad/b_acc_3:0",
_device="/job:localhost/replica:0/task:0/cpu:0"]
(gradients_1/rnn/while/rnn/layer_norm_basic_lstm_cell/state_1/batchnorm/sub/Enter_grad/b_acc_3)]] ""
```

```
In [26]: # !tensorboard --ip 0.0.0.0 --logdir logdir
```

```
In [27]: tf.reset_default_graph()

inputs_ = tf.placeholder(tf.float32,
                          [BATCH_SIZE, MAX_SEQ_LEN, INPUT_UNITS],
                          name='inputs')
labels_ = tf.placeholder(tf.int64,
                          [BATCH_SIZE],
                          name='labels')

model = MnistRnn(inputs_,
                  labels_,
                  INPUT_UNITS,
                  NUM_HIDDEN_UNITS,
                  BATCH_SIZE,
                  MAX_SEQ_LEN,
                  lambda num_hidden_units: \
                      tf.contrib.rnn.LayerNormBasicLSTMCell(
                          num_hidden_units,
                          norm_gain=0.85,
                          norm_shift=0.15),
                  add_check = True,
                  lr = 0.0001,
                  use_grad_clip = True)
```

```
In [28]: config = tf.ConfigProto(gpu_options={'allow_growth':True})
sess = tf.InteractiveSession(config=config)

tf.global_variables_initializer().run()

train_writer = tf.summary.FileWriter(
    'logdir/train_ln_basic_lstm_2',
    graph=tf.get_default_graph())
test_writer  = tf.summary.FileWriter(
    'logdir/test_ln_basic_lstm_2',
    graph=tf.get_default_graph())

train(inputs_, labels_, 10, train_writer, test_writer)
```

```
[trn] ep 1, step 250, loss 2.108183, accu 0.046875, sec/iter 0.104071
[tst] ep 1, step 507, accu 0.421875, sec/iter 0.012975
[trn] ep 2, step 750, loss 1.525936, accu 0.429688, sec/iter 0.104317
[tst] ep 2, step 1014, accu 0.515625, sec/iter 0.009750
[trn] ep 3, step 1250, loss 1.210164, accu 0.539062, sec/iter 0.103209
[tst] ep 3, step 1521, accu 0.601562, sec/iter 0.008972
[trn] ep 4, step 1750, loss 0.963128, accu 0.632812, sec/iter 0.088096
[tst] ep 4, step 2028, accu 0.671875, sec/iter 0.008804
[trn] ep 5, step 2250, loss 0.794114, accu 0.664062, sec/iter 0.087969
[tst] ep 5, step 2535, accu 0.687500, sec/iter 0.010850
[trn] ep 6, step 2750, loss 0.673972, accu 0.726562, sec/iter 0.091234
[tst] ep 6, step 3042, accu 0.726562, sec/iter 0.008794
[trn] ep 7, step 3250, loss 0.574158, accu 0.750000, sec/iter 0.092356
[tst] ep 7, step 3549, accu 0.742188, sec/iter 0.008645
[trn] ep 8, step 3750, loss 0.494440, accu 0.750000, sec/iter 0.088173
[tst] ep 8, step 4056, accu 0.757812, sec/iter 0.008989
[trn] ep 9, step 4250, loss 0.431050, accu 0.757812, sec/iter 0.083751
[tst] ep 9, step 4563, accu 0.789062, sec/iter 0.008605
[trn] ep 10, step 4750, loss 0.378663, accu 0.765625, sec/iter 0.085223
[tst] ep 10, step 5070, accu 0.804688, sec/iter 0.008551
```

## 정리해 봅시다

### 텐서플로우에서 지원하는 RNN Cell 유형

- [tf.contrib.rnn.BasicRNNCell\(\)](#)
- [tf.contrib.rnn.BasicLSTMCell\(\)](#)
- [tf.contrib.rnn.LSTMCell\(\)](#)
- [tf.contrib.rnn.GRUCell\(\)](#)
- [tf.contrib.rnn.LayerNormBasicLSTMCell\(\)](#)

### RNN 구성

- [tf.nn.dynamic\\_rnn\(\)](#)

```

cell          = tf.contrib.rnn.BasicRNNCell(
                    num_hidden_units)
last, states = tf.nn.dynamic_rnn(
                    cell,
                    inputs,
                    sequence_length=sequence_length,
                    dtype=tf.float32)

```

## 실행중 NaN, Inf 등이 발생하면?

1. `tf.check_numerics()`로 문제가 되는 컴퍼넌트 파악
2. `learning_rate` 를 줄여본다
3. 적당한 값으로 `gradient clipping` 해 본다
4. 사용한 컴퍼넌트에 `smoothing` 할 수 있는 파라미터가 있는지 확인하고 적용해 본다

## gradient clipping?

- `tf.clip_by_global_norm`

```

clip_by_global_norm(
    t_list,
    clip_norm,
    use_norm=None,
    name=None
)

```

- `tf.gradients`

```

gradients(
    ys,
    xs,
    grad_ys=None,
    name='gradients',
    colocate_gradients_with_ops=False,
    gate_gradients=False,
    aggregation_method=None
)

```

- `optimizer.apply_gradients`

```

apply_gradients(
    grads_and_vars,
    global_step=None,
    name=None
)

```

- 사용예:

```
tvars_      = tf.trainable_variables()
grads_, _   = tf.clip_by_global_norm(
                tf.gradients(loss,tvars_),
                5.0)
optimize    = tf.train.AdamOptimizer(
                learning_rate=learning_rate) \
                .apply_gradients(zip(grads_, tvars_))
```